

# Naïve Bayesians

Hierarchical Modelling I Intro

## Goal

Developing the Bayesian  
muscle to solve a wide  
range of problems

# Naïve Bayesian Philosophy

**Intuitive (Visual)  
Understanding of the  
Bayesian Reasoning**

**Ability to model real  
world problems in a  
Bayesian Setting**

**Fluency in the Calculus  
of Bayesian Stats & ML  
model**

Starting from Simple  
Probabilistic modelling

Adapting it in a a Bayesian  
setting  
And moving towards ML  
models



# Bayes Rule

Posterior

Likelihood

Prior

$$P(\theta | D) = \frac{P(D | \theta) P(\theta)}{P(D)}$$

Normalising Constant

# Bayes Rule

Posterior

Likelihood

Prior

$$P(\theta_i | D) = \frac{P(D | \theta_i) P(\theta_i)}{\sum_{all\ j} P(D | \theta_j) P(\theta_j)}$$

Normalising Constant

# Bayes Rule

Posterior

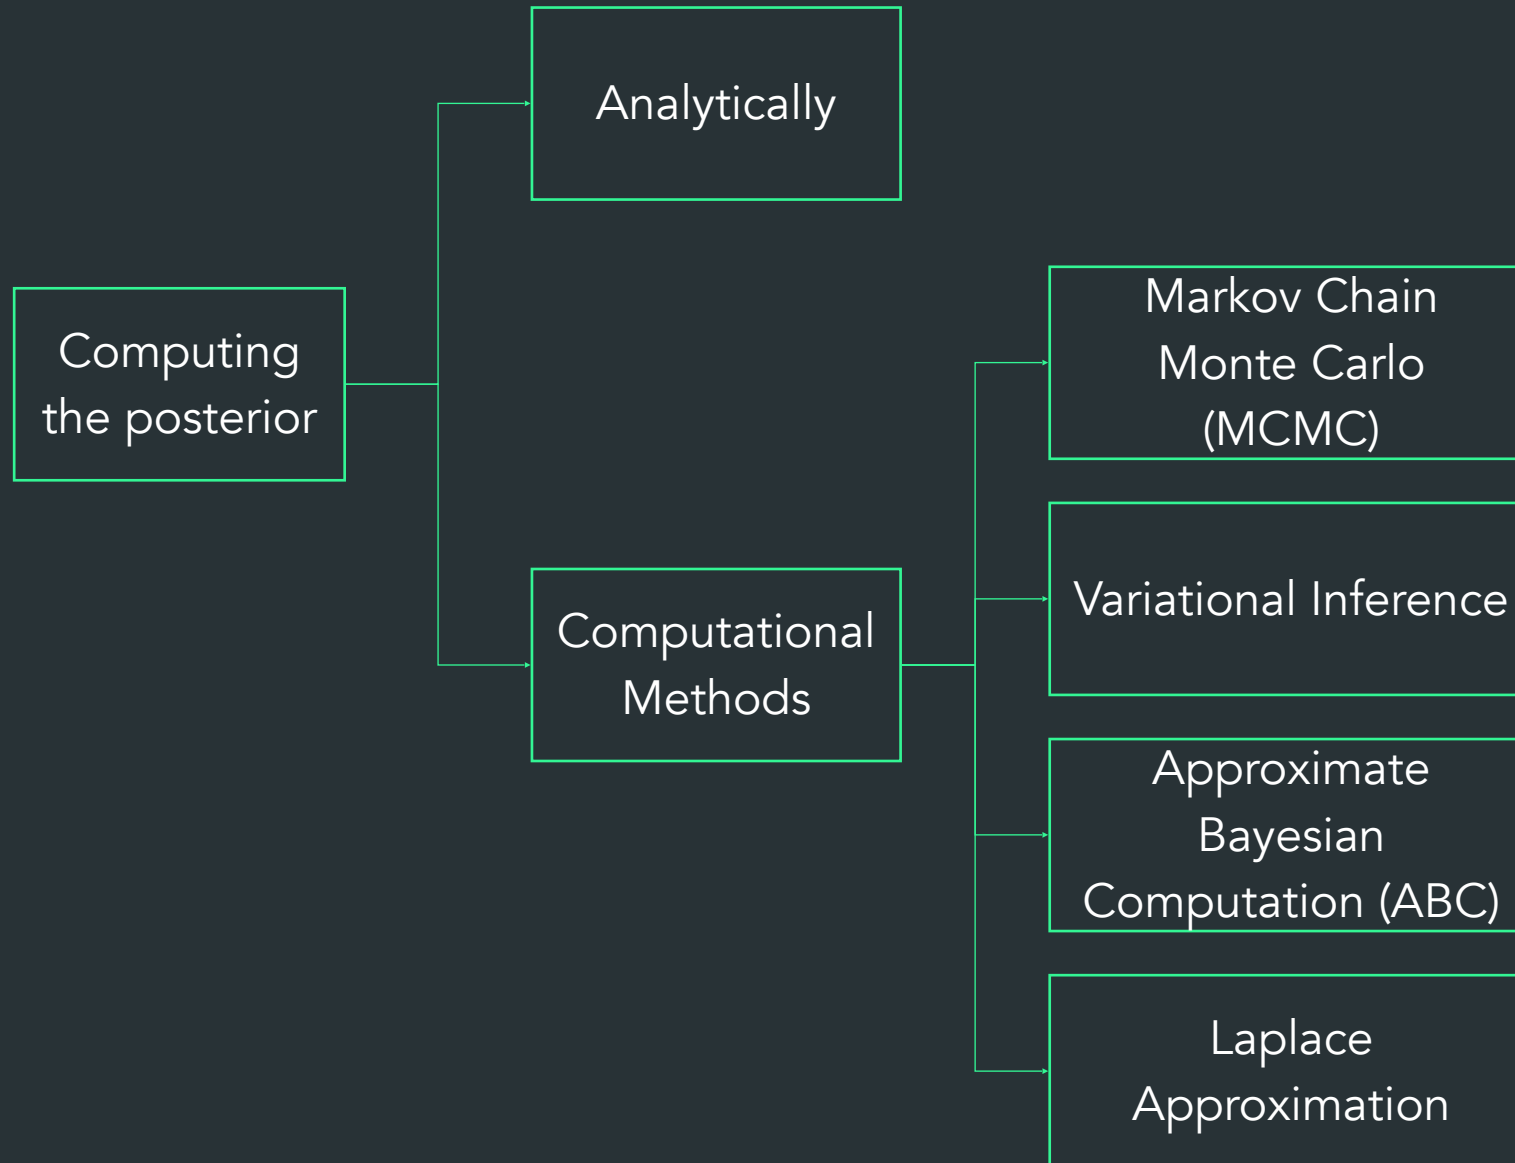
Likelihood

Prior

$$P(\theta | D) = \frac{P(D | \theta) P(\theta)}{\int P(D | \theta') P(\theta') d\theta'}$$

Normalising Constant

# Ways to get to the Posterior



## Generating samples from a random variable

$$x_k \sim \mathcal{N}(\mu = 0, \sigma = 1)$$



# Generating samples from a random variable

$$x_k \sim \mathcal{N}(\mu = 0, \sigma = 1)$$

```
from scipy import stats

# Sample the distribution
x = stats.norm.rvs(
    size=1000,
    loc=0,
    scale=1,
)
```

$$x = \begin{bmatrix} x_1 \\ \vdots \\ x_K \end{bmatrix}$$

# Generating samples from a random variable

$$x_k \sim \mathcal{N}(\mu = 0, \sigma = 1)$$

```
from scipy import stats

# Sample the distribution
x = stats.norm.rvs(
    size=1000,
    loc=0,
    scale=1,
)
```

```
import pymc3 as pm

with pm.Model() as model:
    # Specify the distribution
    x_n = pm.Normal(name="x", mu=0, sigma=1)
    # Sample
    trace = pm.sample(draws=1000)

x_pm = trace.get_values(varname="x")
```

$$x = \begin{bmatrix} x_1 \\ \vdots \\ x_K \end{bmatrix}$$

# Fitting a probability distribution

$$x_k \sim \mathcal{N}(\mu = 0, \sigma = 1)$$

```
from scipy import stats

# Fitting the distribution
mu_fit, sigma_fit = stats.norm.fit(data=x)
```

```
import pymc3 as pm

with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

# Fitting a probability distribution

$$x_k \sim \mathcal{N}(\mu = 0, \sigma = 1)$$

```
from scipy import stats

# Fitting the distribution
mu_fit, sigma_fit = stats.norm.fit(data=x)
```

```
import pymc3 as pm

with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

# Fitting a probability distribution

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

Initialise Model

# Fitting a probability distribution

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

Define priors

$$\mu \sim \mathcal{N}(\mu = 0, \sigma = 10)$$

# Fitting a probability distribution

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

Define likelihood

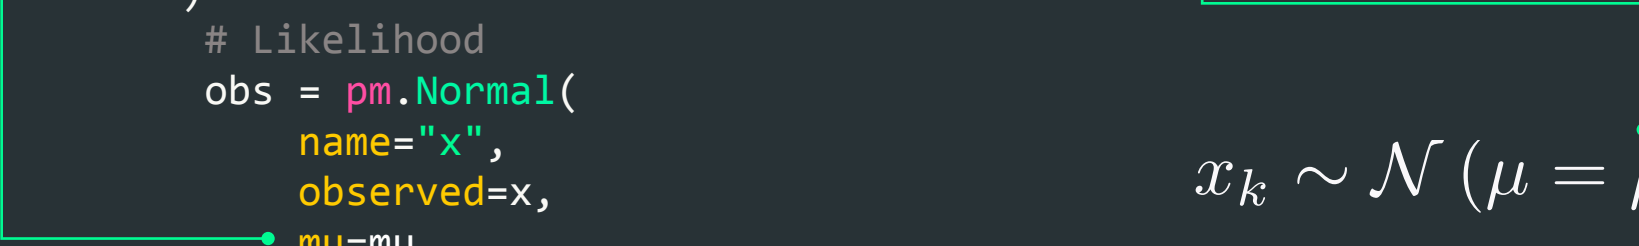
$$x_k \sim \mathcal{N}(\mu = \mu, \sigma = 1)$$

# Fitting a probability distribution

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```


$$\mu \sim \mathcal{N}(\mu = 0, \sigma = 10)$$


$$x_k \sim \mathcal{N}(\mu = \mu, \sigma = 1)$$



# Fitting a probability distribution

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

Pass in the data you have

# Fitting a probability distribution

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

Compute the posterior

# Fitting a probability distribution

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

Get samples from your posterior distribution

# Fitting a probability distribution | Vectorising the parameters

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

# Fitting a probability distribution | Vectorising the parameters

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
        shape=1,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

Vectorise the prior parameters

$$\mu = [\mu \sim \mathcal{N}(\mu = 0, \sigma = 10)]$$

# Fitting a probability distribution | Vectorising the parameters

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
        shape=1,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x,
        mu=mu,
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

Vectorise the prior parameters

$$\mu = [\mu \sim \mathcal{N}(\mu = 0, \sigma = 10)]$$

# Fitting a probability distribution | Vectorising the parameters

```
index = np.zeros(len(x), dtype=np.int8)
```

```
with pm.Model() as mod:
```

```
    # Prior
```

```
    mu = pm.Normal(  
        name="mu",  
        mu=0,  
        sigma=10,  
        shape=1,  
    )
```

```
    # Likelihood
```

```
    obs = pm.Normal(  
        name="x",  
        observed=x,  
        mu=mu,  
        sigma=1,  
    )
```

```
    # Posterior
```

```
    trace = pm.sample(draws=1000)
```

```
mu_pm = trace.get_values(varname="mu")
```

# Fitting a probability distribution | Vectorising the parameters

```
index = np.zeros(len(x), dtype=np.int8)
```

```
with pm.Model() as mod:
```

```
    # Prior
```

```
    mu = pm.Normal(  
        name="mu",  
        mu=0,  
        sigma=10,  
        shape=1,  
    )
```

```
    # Likelihood
```

```
    obs = pm.Normal(  
        name="x",  
        observed=x,  
        mu=mu,  
        sigma=1,  
    )
```

```
    # Posterior
```

```
    trace = pm.sample(draws=1000)
```

```
mu_pm = trace.get_values(varname="mu")
```

$$x = \begin{bmatrix} \mathcal{N}(\mu = \mu, \sigma = 1) \\ \vdots \\ \mathcal{N}(\mu = \mu, \sigma = 1) \end{bmatrix}$$



# Fitting a probability distribution | Vectorising the parameters

```
index = np.zeros(len(x), dtype=np.int8)
```

```
with pm.Model() as mod:
```

```
    # Prior
```

```
    mu = pm.Normal(  
        name="mu",  
        mu=0,  
        sigma=10,  
        shape=1,  
    )
```

```
    # Likelihood
```

```
    obs = pm.Normal(  
        name="x",  
        observed=x,  
        mu=mu,  
        sigma=1,  
    )
```

```
    # Posterior
```

```
    trace = pm.sample(draws=1000)
```

```
mu_pm = trace.get_values(varname="mu")
```

# Fitting a probability distribution | Vectorising the parameters

```
index = np.zeros(len(x), dtype=np.int8)
```

```
with pm.Model() as mod:
```

```
    # Prior
```

```
    mu = pm.Normal(  
        name="mu",  
        mu=0,  
        sigma=10,  
        shape=1,  
    )
```

```
    # Likelihood
```

```
    obs = pm.Normal(  
        name="x",  
        observed=x,  
        mu=mu[index],  
        sigma=1,  
    )
```

```
    # Posterior
```

```
    trace = pm.sample(draws=1000)
```

```
mu_pm = trace.get_values(varname="mu")
```

"Explode" the parameters


$$\begin{bmatrix} \mu \\ \vdots \\ \mu \end{bmatrix}$$

*Length of the data*

# Fitting a probability distribution | Vectorising the parameters

```
index = np.zeros(len(x), dtype=np.int8)
```

```
with pm.Model() as mod:
```

```
    # Prior
```

```
    mu = pm.Normal(  
        name="mu",  
        mu=0,  
        sigma=10,  
        shape=1,  
    )
```

```
    # Likelihood
```

```
    obs = pm.Normal(  
        name="x",  
        observed=x,  
        mu=mu[index],  
        sigma=1,  
    )
```

```
    # Posterior
```

```
    trace = pm.sample(draws=1000)
```

```
mu_pm = trace.get_values(varname="mu")
```

# Moving to hierarchical models

```
# Generate synthetic hierarchical data
x_a = stats.norm.rvs(
    size=1000,
    loc=0,
    scale=1,
)
x_b = stats.norm.rvs(
    size=100,
    loc=2,
    scale=1,
)
x_hierarchy = np.append(x_a, x_b)
```

$$\mathbf{x} = \begin{bmatrix} x_1^a \\ \vdots \\ x_K^a \\ x_1^b \\ \vdots \\ x_{K'}^b \end{bmatrix}$$

# Moving to hierarchical models | Pooled Model

```
# Generate synthetic hierarchical data
x_a = stats.norm.rvs(
    size=1000,
    loc=0,
    scale=1,
)
x_b = stats.norm.rvs(
    size=200,
    loc=2,
    scale=1,
)
x_hierarchy = np.append(x_a, x_b)

# Fitting the distribution
mu_fit, sigma_fit = stats.norm.fit(data=x_hierarchy)
```

$$\mathbf{x} = \begin{bmatrix} x_1^a \\ \vdots \\ x_K^a \\ x_1^b \\ \vdots \\ x_{K'}^b \end{bmatrix}$$

# Moving to hierarchical models | Unpooled Model

```
index_a = np.zeros(len(x_a), dtype=np.int8)
index_b = np.ones(len(x_b), dtype=np.int8)
index = np.append(index_a, index_b)
```

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
        shape=2,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x_hierarchy,
        mu=mu[index],
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

$$\mu = \begin{bmatrix} \mu_a \sim \mathcal{N}(\mu = 0, \sigma = 1) \\ \mu_b \sim \mathcal{N}(\mu = 0, \sigma = 1) \end{bmatrix}$$

# Moving to hierarchical models | Unpooled Model

```
index_a = np.zeros(len(x_a), dtype=np.int8)
index_b = np.ones(len(x_b), dtype=np.int8)
index = np.append(index_a, index_b)
```

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
        shape=2,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x_hierarchy,
        mu=mu[index],
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```

$$i = \begin{bmatrix} 0 \\ \vdots \\ 0 \\ 1 \\ \vdots \\ 1 \end{bmatrix}$$

# Moving to hierarchical models | Unpooled Model

```
index_a = np.zeros(len(x_a), dtype=np.int8)
index_b = np.ones(len(x_b), dtype=np.int8)
index = np.append(index_a, index_b)
```

```
with pm.Model() as mod:
```

```
    # Prior
```

```
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
        shape=2,
    )
```

```
    # Likelihood
```

```
    obs = pm.Normal(
        name="x",
        observed=x_hierarchy,
        mu=mu[index],
        sigma=1,
    )
```

```
    # Posterior
```

```
    trace = pm.sample(draws=1000)
```

```
mu_pm = trace.get_values(varname="mu")
```

$$\bullet \text{-----} \mathbf{x} = \begin{bmatrix} x_1^a \\ \vdots \\ x_K^a \\ x_1^b \\ \vdots \\ x_{K'}^b \end{bmatrix}$$



# Moving to hierarchical models | Unpooled Model

```
index_a = np.zeros(len(x_a), dtype=np.int8)
index_b = np.ones(len(x_b), dtype=np.int8)
index = np.append(index_a, index_b)
```

```
with pm.Model() as mod:
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=10,
        shape=2,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x_hierarchy,
        mu=mu[index],
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)

mu_pm = trace.get_values(varname="mu")
```


$$\begin{bmatrix} \mu_a \\ \vdots \\ \mu_a \\ \mu_b \\ \vdots \\ \mu_b \end{bmatrix}$$

# Moving to hierarchical models | Unpooled Model

```
with pm.Model() as hierarchical_mod:
```

```
    # Prior
    mu = pm.Normal(
        name="mu",
        mu=0,
        sigma=0.1,
        shape=2,
    )
    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x_hierarchy,
        mu=mu[index],
        sigma=1,
    )
    # Posterior
    trace = pm.sample(draws=1000)
```

$$\mu = \begin{bmatrix} \mu_a \sim \mathcal{N}(\mu = 0, \sigma = 1) \\ \mu_b \sim \mathcal{N}(\mu = 0, \sigma = 1) \end{bmatrix}$$

# Moving to hierarchical models

```
with pm.Model() as hierarchical_mod:
    # Set a global prior for the parameters
    mu_glob = pm.Normal(
        name="mu_glob",
        mu=1,
        sigma=0.1,
    )

    # Prior
    mu = pm.Normal(
        name="mu",
        mu=mu_glob,
        sigma=0.1,
        shape=2,
    )

    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x_hierarchy,
        mu=mu[index],
        sigma=1,
    )

    # Posterior
    trace = pm.sample(draws=1000)
```

$$\mu_{\text{glob}} \sim \mathcal{N}(\mu = 1, \sigma = 0.1)$$

# Moving to hierarchical models

```
with pm.Model() as hierarchical_mod:
    # Set a global prior for the parameters
    mu_glob = pm.Normal(
        name="mu_glob",
        mu=1,
        sigma=0.1,
    )

    # Prior
    mu = pm.Normal(
        name="mu",
        mu=mu_glob,
        sigma=0.1,
        shape=2,
    )

    # Likelihood
    obs = pm.Normal(
        name="x",
        observed=x_hierarchy,
        mu=mu[index],
        sigma=1,
    )

    # Posterior
    trace = pm.sample(draws=1000)
```

$$\mu_{\text{glob}} \sim \mathcal{N}(\mu = 1, \sigma = 0.1)$$

$$\mu = \begin{bmatrix} \mu_a \sim \mathcal{N}(\mu = \mu_{\text{glob}}, \sigma = 0.1) \\ \mu_b \sim \mathcal{N}(\mu = \mu_{\text{glob}}, \sigma = 0.1) \end{bmatrix}$$

# Moving to hierarchical models

